

**# Name: Athiena Rachel**

**# Roll no.: 230701016**

**# Date: 30.07.2024**


**#Exercise: Matplotlibseaborn**

```
import pandas as pd
import matplotlib.pyplot as plt

data = {'Year': list(range(2010, 2021)),
        'Job Postings': [150, 300, 450, 600, 800, 1200, 1600, 2100, 2700, 3400, 4200]}

df = pd.DataFrame(data)

plt.plot(df['Year'], df['Job Postings'], marker='o')
plt.title('Trend of Data Science Job Postings')
plt.xlabel('Year')
plt.ylabel('Number of Job Postings')
plt.show()
```



```
import pandas as pd

import matplotlib.pyplot as plt

roles=['data scientist','data analyst','data engineer','ml engineer','business analyst']

counts=[300,500,450,200,150]

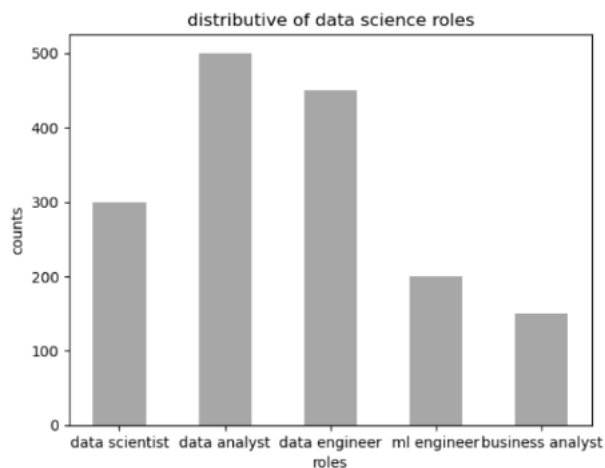
plt.bar(roles,counts,width=0.5,color='yellow')

plt.title('distributive of data science roles')

plt.xlabel('roles')

plt.ylabel('counts')

plt.show()
```



```
import pandas as pd

structured_data = pd. DataFrame ({'ID's [1, 2, 3], 'NAME': ['Aish', 'Betty', 'Cathy'], 'AGE': [18, 20, 25], 'GRADE': ['O', 'A', 'B'], 'SKILL': ['Art', 'Music', 'Dance']})

print ("Structured Data: \n", structured_data)

unstructured_data = "This is an unstructured data. It can be text, an image, or a video"

print("unstructured data: \n", unstructured_data)

semistricted_clata = {'ID': 1, 'NAME': 'Alice', 'ATTRIBUTES':
```

```
'HEIGHT': 170, 'WEIGHT: 45}}  
print("Semistructured Data: \n", semistructured_data)
```

## OUTPUT-

### Structured Data:

```
'ID's [1, 2, 3], 'NAME': ['Aish', 'Betty', 'Cathy'], 'AGE': [18, 20, 25], 'GRADE': ['O',  
'A', 'B'],  
'SKILL': ['Art', 'Music', 'Dance']
```

### unstructured data:

**This is an unstructured data. It can be text, an image, or a video**

### Semistructured Data:

```
'ID': 1, 'NAME': 'Alice', 'ATTRIBUTES':  
'HEIGHT': 170, 'WEIGHT: 45}
```

```
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
df=pd.read_csv ('sales_data.csv');  
print(df.head());  
print (df.isnull().sum())  
df ['Sales ']. fillna(df ['Sales']. mean(), inplace = True)  
df. dropna (subset = ['Product', 'Quantity', 'Region'], inplace = True)  
print (df. describe())  
prodsummary= df.groupby ('Product'). agg ({'sales': 'sum', 'Quantity': 'sum'}).reset_index()  
print (prodsummary)
```

---

	Date	Product	Sales	Quantity	Region
0	01-01-2023	Product A	200	4	North
1	02-01-2023	Product B	150	3	South
2	03-01-2023	Product A	220	5	North
3	04-01-2023	Product C	300	6	East
4	05-01-2023	Product B	180	4	West

Date 0

Product 0

Sales 0

Quantity 0

Region 0

dtype: int64

---

**# Name: Athiena Rachel**

**# Roll no.: 230701016**

**# Date: 06.08.2024**

**#Exercise: Pandas Built in function; Numpy Built in function- Array slicing, Ravel, Reshape, ndim**

```
import numpy as np
array = np.random.randint(1, 100, 9)
np.sqrt(array)
array.ndim
new_array = array.reshape(3, 3)
new_array.ndim
new_array.ravel()
newm = new_array.reshape(3, 3)
newm[2, 1:3]
newm[1:2, 1:3]
new_array[0:3, 0:0]
new_array[0:2, 0:1]
new_array[0:3, 0:1]
new_array[1:3]
```

```
newm[2,1:3]
↳ array([39, 51])

newm[1:2,1:3]
↳ array([[62, 15]])

new_array[0:3,0:0]
↳ array([], shape=(3, 0), dtype=int64)

new_array[0:2,0:1]
↳ array([[83],
        [47]])

new_array[0:3,0:1]
↳ array([[83],
        [47],
        [96]])

new_array[1:3]
↳ array([[47, 62, 15],
        [96, 39, 51]])
```

**# Name: Athiena Rachel**

**# Roll no.: 230701016**

**# Date: 28.08.2024**

**#Exercise: Outlier detection**

```
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
array = np.random.randint(1, 100, 16)
print(array)
print(array.mean())
print(np.percentile(array, 25))
print(np.percentile(array, 50))
print(np.percentile(array, 75))
print(np.percentile(array, 100))
def outDetection(array):
    sorted_array = sorted(array)
    Q1, Q3 = np.percentile(array, [25, 75])
    IQR = Q3 - Q1

    lr = Q1 - (1.5 * IQR)
    ur = Q3 + (1.5 * IQR)

    return lr, ur

lr, ur = outDetection(array)
print(f"Lower Range: {lr}, Upper Range: {ur}")

sns.displot(array)
```


```
plt.show()
```

```
new_array = array[(array > lr) & (array < ur)]  
print("Array after outlier removal:", new_array)
```



```
sns.displot(new_array)
```


```
plt.show()
```

```
lr1, ur1 = outDetection(new_array)  
print(f"New Lower Range: {lr1}, New Upper Range: {ur1}")  
final_array = new_array[(new_array > lr1) & (new_array < ur1)]  
print("Final array after second outlier removal:", final_array)  
sns.displot(final_array)  
plt.show()
```









**# Name: Athiena Rachel**

**# Roll no.: 230701016**

**# Date: 27.08.2024**

**#Exercise: Missing and inappropriate data**

```
import numpy as np
```

```
import pandas as pd
```

```
df = pd.read_csv("Hotel_Dataset.csv")
```

```
df
```

```
df.duplicated()
```

```
df.info()
```

```
df.drop_duplicates(inplace=True)
```

```
df
```

```
len(df)
```

```
index = np.array(list(range(0, len(df))))
```

```
df.set_index(index, inplace=True)
```

```
index
```

```
df.drop(['Age_Group.1'], axis=1, inplace=True)
```

```
df
```

```
df.CustomerID.loc[df.CustomerID < 0] = np.nan
```

```
df.Bill.loc[df.Bill < 0] = np.nan
```

```
df.EstimatedSalary.loc[df.EstimatedSalary < 0] = np.nan
```

```
df
```

```
df['NoOfPax'].loc[(df['NoOfPax'] < 1) | (df['NoOfPax'] > 20)] = np.nan
```

```
df
```

```
df.Age_Group.unique()
```

```
df.Hotel.unique()
```

```
df.Hotel.replace(['Ibys'], 'Ibis', inplace=True)
```

```
df.FoodPreference.unique()
```

```
df.FoodPreference.replace(['Vegetarian', 'veg'], 'Veg', inplace=True)
```

```
df.FoodPreference.replace(['non-Veg'], 'Non-Veg', inplace=True)
```

```
df.EstimatedSalary.fillna(round(df.EstimatedSalary.mean()), inplace=True)
```

```
df.NoOfPax.fillna(round(df.NoOfPax.median()), inplace=True)
```

```
df['Rating(1-5)'].fillna(round(df['Rating(1-5)'].median()), inplace=True)
```

```
df.Bill.fillna(round(df.Bill.mean()), inplace=True)
```

```
df
```

	CustomerID	Age_Group	Rating(1-5)	Hotel	FoodPreference	Bill	NoOfPax	EstimatedSalary	Age_Group.1
0	1	20-25	4	Ibis	veg	1300	2	40000	20-25
1	2	30-35	5	LemonTree	Non-Veg	2000	3	59000	30-35
2	3	25-30	6	RedFox	Veg	1322	2	30000	25-30
3	4	20-25	-1	LemonTree	Veg	1234	2	120000	20-25
4	5	35+	3	Ibis	Vegetarian	989	2	45000	35+
5	6	35+	3	Ibys	Non-Veg	1909	2	122220	35+
6	7	35+	4	RedFox	Vegetarian	1000	-1	21122	35+
7	8	20-25	7	LemonTree	Veg	2999	-10	345673	20-25
8	9	25-30	2	Ibis	Non-Veg	3456	3	-99999	25-30
9	9	25-30	2	Ibis	Non-Veg	3456	3	-99999	25-30
10	10	30-35	5	RedFox	non-Veg	-6755	4	87777	30-35

```

0      False
1      False
2      False
3      False
4      False
5      False
6      False
7      False
8      False
9       True
10     False
dtype: bool

```

---

	CustomerID	Age_Group	Rating(1-5)	Hotel	FoodPreference	Bill	NoOfPax	EstimatedSalary	Age_Group.1
0	1	20-25	4	Ibis	veg	1300	2	40000	20-25
1	2	30-35	5	LemonTree	Non-Veg	2000	3	59000	30-35
2	3	25-30	6	RedFox	Veg	1322	2	30000	25-30
3	4	20-25	-1	LemonTree	Veg	1234	2	120000	20-25
4	5	35+	3	Ibis	Vegetarian	989	2	45000	35+
5	6	35+	3	Ibys	Non-Veg	1909	2	122220	35+
6	7	35+	4	RedFox	Vegetarian	1000	-1	21122	35+
7	8	20-25	7	LemonTree	Veg	2999	-10	345673	20-25
8	9	25-30	2	Ibis	Non-Veg	3456	3	-99999	25-30
10	10	30-35	5	RedFox	non-Veg	-6755	4	87777	30-35

	CustomerID	Age_Group	Rating(1-5)	Hotel	FoodPreference	Bill	NoOfPax	EstimatedSalary
0	1	20-25	4	Ibis	veg	1300	2	40000
1	2	30-35	5	LemonTree	Non-Veg	2000	3	59000
2	3	25-30	6	RedFox	Veg	1322	2	30000
3	4	20-25	-1	LemonTree	Veg	1234	2	120000
4	5	35+	3	Ibis	Vegetarian	989	2	45000
5	6	35+	3	Ibys	Non-Veg	1909	2	122220
6	7	35+	4	RedFox	Vegetarian	1000	-1	21122
7	8	20-25	7	LemonTree	Veg	2999	-10	345673
8	9	25-30	2	Ibis	Non-Veg	3456	3	-99999
9	10	30-35	5	RedFox	non-Veg	-6755	4	87777

	CustomerID	Age_Group	Rating(1-5)	Hotel	FoodPreference	Bill	NoOfPax	EstimatedSalary
0	1.0	20-25	4.0	Ibis	Veg	1300.0	2.0	40000.0
1	2.0	30-35	5.0	LemonTree	Non-Veg	2000.0	3.0	59000.0
2	3.0	25-30	4.0	RedFox	Veg	1322.0	2.0	30000.0
3	4.0	20-25	4.0	LemonTree	Veg	1234.0	2.0	120000.0
4	5.0	35+	3.0	Ibis	Veg	989.0	2.0	45000.0
5	6.0	35+	3.0	Ibis	Non-Veg	1909.0	2.0	122220.0
6	7.0	35+	4.0	RedFox	Veg	1000.0	2.0	21122.0
7	8.0	20-25	4.0	LemonTree	Veg	2999.0	2.0	345673.0
8	9.0	25-30	2.0	Ibis	Non-Veg	3456.0	3.0	96755.0
9	10.0	30-35	5.0	RedFox	Non-Veg	1801.0	4.0	87777.0

**# Name: Athiena Rachel**

**# Roll no.: 230701016**

**# Date: 30.08.2024**

**#Exercise: Data Preprocessing**

```
import numpy as np
```

```
import pandas as pd
```

```
df = pd.read_csv("/content/pre-process_datasample.csv")
```

```
df.info()
```

```
df.Country.fillna(df.Country.mode()[0], inplace=True)
```

```
df.Age.fillna(df.Age.median(), inplace=True)
```

```
df.Salary.fillna(round(df.Salary.mean()), inplace=True)
```

```
encoded_countries = pd.get_dummies(df.Country)
```

```
updated_dataset = pd.concat([encoded_countries, df.iloc[:, [1, 2, 3]]], axis=1)
```

```
updated_dataset.Purchased.replace(['No', 'Yes'], [0, 1], inplace=True)
```

```
print(updated_dataset)
```

	France	Germany	Spain	Age	Salary	Purchased
0	True	False	False	44.0	72000.0	0
1	False	False	True	27.0	48000.0	1
2	False	True	False	30.0	54000.0	0
3	False	False	True	38.0	61000.0	0
4	False	True	False	40.0	63778.0	1
5	True	False	False	35.0	58000.0	1
6	False	False	True	38.0	52000.0	0
7	True	False	False	48.0	79000.0	1
8	True	False	False	50.0	83000.0	0
9	True	False	False	37.0	67000.0	1



**# Name: Athiena Rachel**

**# Roll no.: 230701016**

**# Date: 03.09.2024**

**#Exercise: EDA-Quantitative and Qualitative plots**

```
import seaborn as sns
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline

tips = sns.load_dataset('tips')
tips.head()

sns.displot(tips.total_bill, kde=True)
plt.show()

sns.displot(tips.total_bill, kde=False)
plt.show()

sns.jointplot(x=tips.tip, y=tips.total_bill)
plt.show()

sns.jointplot(x=tips.tip, y=tips.total_bill, kind="reg")
plt.show()

sns.jointplot(x=tips.tip, y=tips.total_bill, kind="hex")
plt.show()
```

```
sns.pairplot(tips)
```

```
plt.show()
```

```
print(tips.time.value_counts())
```

```
sns.pairplot(tips, hue='time')
```

```
plt.show()
```

```
sns.pairplot(tips, hue='day')
```

```
plt.show()
```

```
sns.heatmap(tips.corr(numeric_only=True), annot=True)
```

```
plt.show()
```

```
sns.boxplot(x=tips.total_bill)
```

```
plt.show()
```

```
sns.boxplot(x=tips.tip)
```

```
plt.show()
```

```
sns.countplot(x=tips.day)
```

```
plt.show()
```

```
sns.countplot(x=tips.sex)
```

```
plt.show()
```

```
tips.sex.value_counts().plot(kind='pie', autopct='% 1.1f%%')
```



```
plt.show()
```



```
tips.sex.value_counts().plot(kind='bar')
```

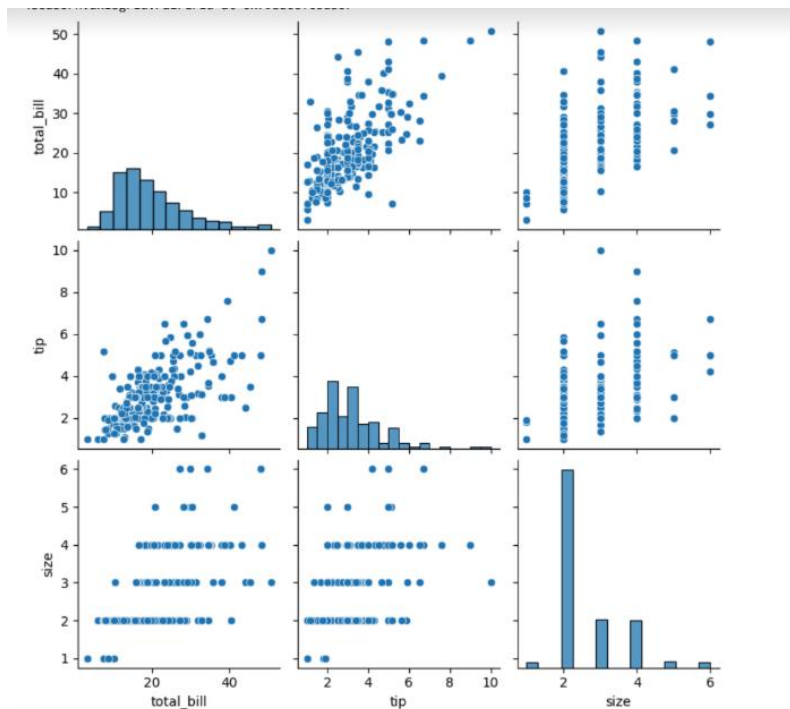

```
plt.show()
```

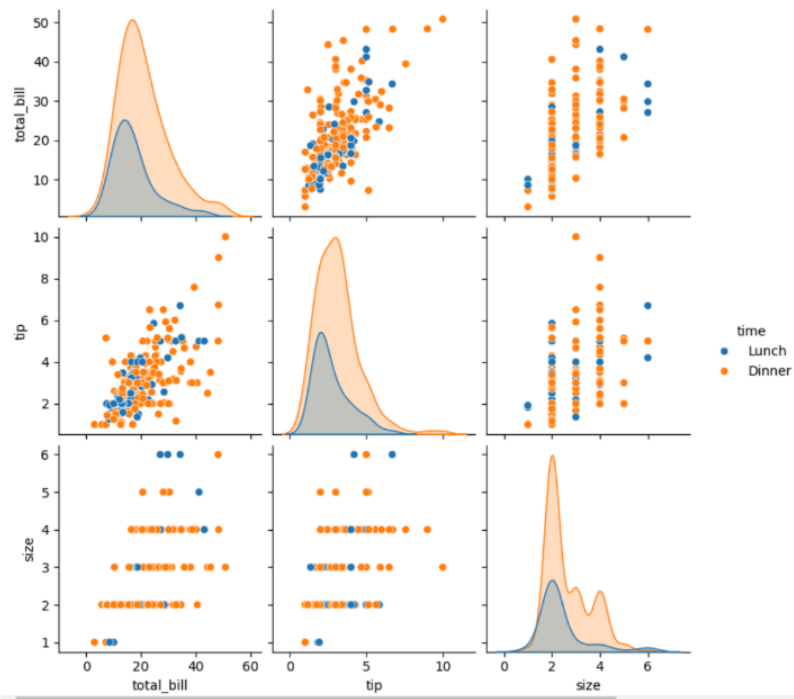
```
sns.countplot(x=tips[tips.time == 'Dinner']['day'])
```

plt.show()









**# Name: Athiena Rachel**

**# Roll no.: 230701016**

**# Date: 10.09.2024**

**#Exercise: Random Sampling and Sampling Distribution**

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
population_mean = 50
```

```
population_std = 10
```

```
population_size = 100000
```

```
population = np.random.normal(population_mean, population_std, population_size)
```

```
plt.figure(figsize=(8, 5))
```

```
plt.hist(population, bins=50, color='skyblue', edgecolor='black', alpha=0.7)
```

```
plt.title('Population Distribution')
```

```
plt.xlabel('Value')
```

```
plt.ylabel('Frequency')
```

```
plt.axvline(population_mean, color='red', linestyle='dashed', linewidth=1.5, label='Population Mean')
```

```
plt.legend()
```

```
plt.show()
```

```
sample_sizes = [30, 50, 100]
```

```
num_samples = 1000
```

```
sample_means = { }
```

```
for size in sample_sizes:
```

```
    sample_means[size] = []
```

```
for _ in range(num_samples):
    sample = np.random.choice(population, size=size, replace=False)
    sample_means[size].append(np.mean(sample))

plt.figure(figsize=(12, 8))
for i, size in enumerate(sample_sizes):
    plt.subplot(len(sample_sizes), 1, i + 1)
    plt.hist(sample_means[size], bins=30, alpha=0.7, color='orange', edgecolor='black',
             label=f'Sample Size {size}')
    plt.axvline(np.mean(population), color='red', linestyle='dashed', linewidth=1.5,
                label='Population Mean')
    plt.title(f'Sampling Distribution of the Sample Mean (Sample Size {size})')
    plt.xlabel('Sample Mean')
    plt.ylabel('Frequency')
    plt.legend()

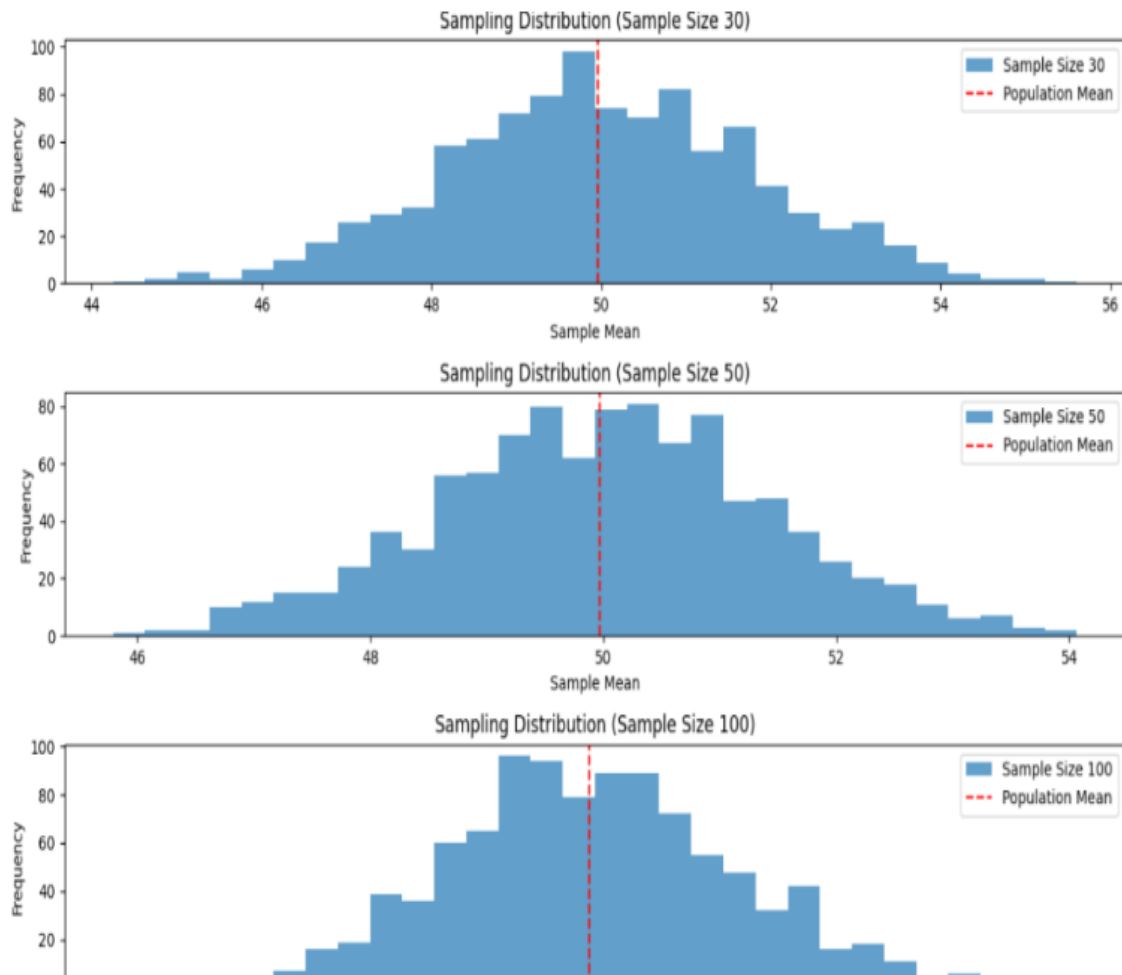
plt.tight_layout()
plt.show()

plt.figure(figsize=(12, 8))
for i, size in enumerate(sample_sizes):
    plt.subplot(len(sample_sizes), 1, i + 1)
    plt.hist(sample_means[size], bins=30, alpha=0.7, color='purple', edgecolor='black',
             label=f'Sample Size {size}', density=True)
    plt.axvline(np.mean(population), color='red', linestyle='dashed', linewidth=1.5,
                label='Population Mean')
    plt.title(f'Sampling Distribution (Sample Size {size}) - CLT Demonstration')
    plt.xlabel('Sample Mean')
    plt.ylabel('Density')
    plt.legend()
```



```
plt.tight_layout()
```

```
plt.show()
```



**# Name: Athiena Rachel**

**# Roll no.: 230701016**

**# Date: 10.09.2024**

**#Exercise: Z-Test**

```
import numpy as np
```

```
import scipy.stats as stats
```

```
sample_data = np.array([
    152, 148, 151, 149, 147, 153, 150, 148, 152, 149,
    151, 150, 149, 152, 151, 148, 150, 152, 149, 150,
    148, 153, 151, 150, 149, 152, 148, 151, 150, 153
])
```

```
population_mean = 150
```

```
sample_mean = np.mean(sample_data)
```

```
sample_std = np.std(sample_data, ddof=1)
```

```
n = len(sample_data)
```

```
z_statistic = (sample_mean - population_mean) / (sample_std / np.sqrt(n))
```

```
p_value = 2 * (1 - stats.norm.cdf(np.abs(z_statistic)))
```

```
print(f"Sample Mean: {sample_mean:.2f}")
```

```
print(f"Z-Statistic: {z_statistic:.4f}")
```

```
print(f"P-Value: {p_value:.4f}")
```

```
alpha = 0.05
```

```
if p_value < alpha:
```

```
    print("Reject the null hypothesis: The average weight is significantly different from 150 grams.")
```

```
else:
```

```
print("Fail to reject the null hypothesis: There is no significant difference in average weight from 150 grams.")
```

**OUTPUT:****Sample Mean: 150.20****Z-Statistic: 0.6406****P-Value: 0.5218****Fail to reject the null hypothesis: There is no significant difference in average weight from 150 grams.**

**# Name: Athiena Rachel**

**# Roll no.: 230701016**

**# Date: 08.10.2024**

**#Exercise: T-Test**

```
import numpy as np
```

```
import scipy.stats as stats
```

```
np.random.seed(42)
```

```
sample_size = 25
```

```
sample_data = np.random.normal(loc=102, scale=15, size=sample_size)
```

```
population_mean = 100
```

```
sample_mean = np.mean(sample_data)
```

```
sample_std = np.std(sample_data, ddof=1)
```

```
n = len(sample_data)
```

```
t_statistic, p_value = stats.ttest_1samp(sample_data, population_mean)
```

```
print(f"Sample Mean: {sample_mean:.2f}")
```

```
print(f"T-Statistic: {t_statistic:.4f}")
```

```
print(f"P-Value: {p_value:.4f}")
```

```
alpha = 0.05
```

```
if p_value < alpha:
```

```
    print("Reject the null hypothesis: The average IQ score is significantly different from  
    100.")
```

else:

```
print("Fail to reject the null hypothesis: There is no significant difference in average IQ score from 100.")
```

### **OUTPUT:**

**Sample Mean: 99.55**

**T-Statistic: -0.1577**

**P-Value: 0.8760**

**Fail to reject the null hypothesis: There is no significant difference in average IQ score from 100.**

**# Name: Athiena Rachel**

**# Roll no.: 230701016**

**# Date: 08.10.2024**

**#Exercise: ANOVA Test**

```
import numpy as np
import scipy.stats as stats
np.random.seed(42)
n_plants = 25
growth_A = np.random.normal(loc=10, scale=2, size=n_plants)
growth_B = np.random.normal(loc=12, scale=3, size=n_plants)
growth_C = np.random.normal(loc=15, scale=2.5, size=n_plants)
f_statistic, p_value = stats.f_oneway(growth_A, growth_B, growth_C)
print("Treatment A Mean Growth:", np.mean(growth_A))
print("Treatment B Mean Growth:", np.mean(growth_B))
print("Treatment C Mean Growth:", np.mean(growth_C))
print()
print(f"F-Statistic: {f_statistic:.4f}")
print(f"P-Value: {p_value:.4f}")
alpha = 0.05
if p_value < alpha:
    print("Reject the null hypothesis: There is a significant difference in mean growth rates
among the three treatments.")
else:
    print("Fail to reject the null hypothesis: There is no significant difference in mean growth
rates among the three treatments.")
if p_value < alpha:
    all_data = np.concatenate([growth_A, growth_B, growth_C])
    treatment_labels = ['A'] * n_plants + ['B'] * n_plants + ['C'] * n_plants

    tukey_results = pairwise_tukeyhsd(all_data, treatment_labels, alpha=0.05)
```

```
print("\nTukey's HSD Post-hoc Test:")  
print(tukey_results)
```

**OUTPUT:****Treatment A Mean Growth: 9.672983882683818****Treatment B Mean Growth: 11.137680744437432****Treatment C Mean Growth: 15.265234904828972****F-Statistic: 36.1214****P-Value: 0.0000****Reject the null hypothesis: There is a significant difference in mean growth rates among the three treatments.**

**# Name: Athiena Rachel**

**# Roll no.: 230701016**

**# Date: 22.10.2024**

**#Exercise: Feature Scaling**

```
import numpy as np
```

```
import pandas as pd
```

```
df = pd.read_csv('/content/pre-process_datasample.csv')
```

```
print("Original Data:")
```

```
print(df)
```

```
df['Country'].fillna(df['Country'].mode()[0], inplace=True)
```

```
features = df.iloc[:, :-1].values
```

```
label = df.iloc[:, -1].values
```

```
from sklearn.impute import SimpleImputer
```

```
age_imputer = SimpleImputer(strategy="mean")
```

```
salary_imputer = SimpleImputer(strategy="mean")
```

```
age_imputer.fit(features[:, [1]])
```

```
salary_imputer.fit(features[:, [2]])
```

```
features[:, [1]] = age_imputer.transform(features[:, [1]])
```

```
features[:, [2]] = salary_imputer.transform(features[:, [2]])
```

```
print("Features after handling missing values:")
```



```
print(features)
```

```
from sklearn.preprocessing import OneHotEncoder
```

```
oh = OneHotEncoder(sparse_output=False)
```

```
Country = oh.fit_transform(features[:, [0]])
```

```
print("OneHotEncoded 'Country' column:")
```

```
print(Country)
```

```
final_set = np.concatenate((Country, features[:, [1, 2]]), axis=1)
```

```
print("Final dataset with OneHotEncoded 'Country' and other features:")
```

```
print(final_set)
```

```
from sklearn.preprocessing import StandardScaler
```

```
sc = StandardScaler()
```

```
sc.fit(final_set)
```

```
feat_standard_scaler = sc.transform(final_set)
```

```
print("Standardized features:")
```

```
print(feat_standard_scaler)
```

```
from sklearn.preprocessing import MinMaxScaler
```

```
mms = MinMaxScaler(feature_range=(0, 1))
```

```
mms.fit(final_set)
```

```
feat_minmax_scaler = mms.transform(final_set)
```

```
print("Normalized features:")
```

```
print(feat_minmax_scaler)
```

```
array([[ 'France', 44.0, 72000.0],
       [ 'Spain', 27.0, 48000.0],
       [ 'Germany', 30.0, 54000.0],
       [ 'Spain', 38.0, 61000.0],
       [ 'Germany', 40.0, 63777.77777777778],
       [ 'France', 35.0, 58000.0],
       [ 'Spain', 38.77777777777778, 52000.0],
       [ 'France', 48.0, 79000.0],
       [ 'France', 50.0, 83000.0],
       [ 'France', 37.0, 67000.0]], dtype=object)
```

```
array([[1.      , 0.      , 0.      , 0.73913043, 0.68571429],
       [0.      , 0.      , 1.      , 0.      , 0.      ],
       [0.      , 1.      , 0.      , 0.13043478, 0.17142857],
       [0.      , 0.      , 1.      , 0.47826087, 0.37142857],
       [0.      , 1.      , 0.      , 0.56521739, 0.45079365],
       [1.      , 0.      , 0.      , 0.34782609, 0.28571429],
       [0.      , 0.      , 1.      , 0.51207729, 0.11428571],
       [1.      , 0.      , 0.      , 0.91304348, 0.88571429],
       [1.      , 0.      , 0.      , 1.      , 1.      ],
       [1.      , 0.      , 0.      , 0.43478261, 0.54285714]])
```

```
array([[1.0, 0.0, 0.0, 44.0, 72000.0],
       [0.0, 0.0, 1.0, 27.0, 48000.0],
       [0.0, 1.0, 0.0, 30.0, 54000.0],
       [0.0, 0.0, 1.0, 38.0, 61000.0],
       [0.0, 1.0, 0.0, 40.0, 63777.77777777778],
       [1.0, 0.0, 0.0, 35.0, 58000.0],
       [0.0, 0.0, 1.0, 38.77777777777778, 52000.0],
       [1.0, 0.0, 0.0, 48.0, 79000.0],
       [1.0, 0.0, 0.0, 50.0, 83000.0],
       [1.0, 0.0, 0.0, 37.0, 67000.0]], dtype=object)
```

```
array([[1.      , 0.      , 0.      , 0.73913043, 0.68571429],
       [0.      , 0.      , 1.      , 0.      , 0.      ],
       [0.      , 1.      , 0.      , 0.13043478, 0.17142857],
       [0.      , 0.      , 1.      , 0.47826087, 0.37142857],
       [0.      , 1.      , 0.      , 0.56521739, 0.45079365],
       [1.      , 0.      , 0.      , 0.34782609, 0.28571429],
       [0.      , 0.      , 1.      , 0.51207729, 0.11428571],
       [1.      , 0.      , 0.      , 0.91304348, 0.88571429],
       [1.      , 0.      , 0.      , 1.      , 1.      ],
       [1.      , 0.      , 0.      , 0.43478261, 0.54285714]])
```

```
array([[ 1.00000000e+00, -5.00000000e-01, -6.54653671e-01,
        7.58874362e-01,  7.49473254e-01],
       [-1.00000000e+00, -5.00000000e-01,  1.52752523e+00,
        -1.71150388e+00, -1.43817841e+00],
       [-1.00000000e+00,  2.00000000e+00, -6.54653671e-01,
        -1.27555478e+00, -8.91265492e-01],
       [-1.00000000e+00, -5.00000000e-01,  1.52752523e+00,
        -1.13023841e-01, -2.53200424e-01],
       [-1.00000000e+00,  2.00000000e+00, -6.54653671e-01,
        1.77608893e-01,  6.63219199e-16],
       [ 1.00000000e+00, -5.00000000e-01, -6.54653671e-01,
        -5.48972942e-01, -5.26656882e-01],
       [-1.00000000e+00, -5.00000000e-01,  1.52752523e+00,
        0.00000000e+00, -1.07356980e+00],
       [ 1.00000000e+00, -5.00000000e-01, -6.54653671e-01,
        1.34013983e+00,  1.38753832e+00],
       [ 1.00000000e+00, -5.00000000e-01, -6.54653671e-01,
        1.63077256e+00,  1.75214693e+00],
       [ 1.00000000e+00, -5.00000000e-01, -6.54653671e-01,
        -2.58340208e-01,  2.93712492e-01]])
```

**# Name: Athiena Rachel**

**# Roll no.: 230701016**

**# Date: 29.10.2024**

**#Exercise: Linear Regression**

```
import numpy as np
```

```
import pandas as pd
```

```
df=pd.read_csv('Salary_data.csv')
```

```
df
```

```
df.info()
```

```
df.dropna(inplace=True)
```

```
df.info()
```

```
df.describe()
```

```
features=df.iloc[:,[0]].values
```

```
label=df.iloc[:,[1]].values
```

```
from sklearn.model_selection import train_test_split
```

```
x_train,x_test,y_train,y_test=train_test_split(features,label,test_size=0.2,random_state=23)
```

```
from sklearn.linear_model import LinearRegression
```

```
model=LinearRegression()
```

```
model.fit(x_train,y_train)
```

```
model.score(x_train,y_train)
```

```
model.score(x_test,y_test)
```

```
model.coef_
```

```
model.intercept_
```

```
import pickle
```

```
pickle.dump(model,open('SalaryPred.model','wb'))
```

```
model=pickle.load(open('SalaryPred.model','rb'))
```

```
yr_of_exp=float(input("Enter Years of Experience: "))
```

```
yr_of_exp_NP=np.array([[yr_of_exp]])
```

```
Salary=model.predict(yr_of_exp_NP)
```

```
print("Estimated Salary for { } years of experience is { }:".format(yr_of_exp,Salary))
```

### **OUTPUT-**

**Estimated Salary for 44.0 years of experience is [[435544.30953887]]:**

**# Name: Athiena Rachel**

**# Roll no.:05.11.2024**

**# Date: 29.10.2024**

**#Exercise: Logistic Regression**

```
import numpy as np
```

```
import pandas as pd
```

```
df=pd.read_csv('Social_Network_Ads.csv')
```

```
df
```

```
df.head()
```

```
features=df.iloc[:,[2,3]].values
```

```
label=df.iloc[:,4].values
```

```
features
```

```
label
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.linear_model import LogisticRegression
```

```
for i in range(1,401):
```

```
    x_train,x_test,y_train,y_test=train_test_split(features,label,test_size=0.2,random_state=i)
```

```
    model=LogisticRegression()
```

```
    model.fit(x_train,y_train)
```

```
    train_score=model.score(x_train,y_train)
```

```
    test_score=model.score(x_test,y_test)
```

```
    if test_score>train_score:
```

```
        print("Test { } Train{ } Random State { }".format(test_score,train_score,i))
```

```
x_train,x_test,y_train,y_test=train_test_split(features,label,test_size=0.2,random_state=314)
```

```
finalModel=LogisticRegression()
```

```
finalModel.fit(x_train,y_train)
```

```
print(finalModel.score(x_train,y_train))
```

```
print(finalModel.score(x_test,y_test))
```

```
from sklearn.metrics import classification_report
```

```
print(classification_report(label,finalModel.predict(features)))
```

	precision	recall	f1-score	support
0	0.85	0.93	0.89	257
1	0.84	0.71	0.77	143
accuracy			0.85	400
macro avg	0.85	0.82	0.83	400
weighted avg	0.85	0.85	0.85	400